

# perf

(Linux profiling with performance counter)

---

Concurrent Programming

# Introduction

---

- What is Perf?
- Installing Perf
- Perf Commands
- Example

# What is Perf?

- Performance analyzing tool in Linux
  - available from Linux kernel version 2.6.31
- Capable of statistical profiling of the entire system
  - (both kernel and userland code)
- Available to measure many types of event
  - (Hardware events, Software events, ...)

# Perf Commands

- perf list  
Show the list of available events to be measured with perf
- perf stat  
Obtain aggregated event counts
- perf record  
Sampling events on per-thread, per-process, per-cpu basis and generate the output file
- perf report  
Analyze the output file generated from *perf record*
- perf annotate  
Analyze the output file generated from *perf record* in the instruction level
- more Commands, but not today

# Perf Commands – perf list

- Show the list of events that can be measured by perf

```
$ perf list
```

```
jongbin@multicore-96:~/TA/Multicore/lab02$ perf list
List of pre-defined events (to be used in -e):

branch-instructions OR branches      [Hardware event]
branch-misses                        [Hardware event]
bus-cycles                           [Hardware event]
cache-misses                         [Hardware event]
cache-references                     [Hardware event]
cpu-cycles OR cycles                 [Hardware event]
instructions                         [Hardware event]
ref-cycles                           [Hardware event]

alignment-faults                     [Software event]
bpf-output                           [Software event]
context-switches OR cs               [Software event]
cpu-clock                            [Software event]
cpu-migrations OR migrations         [Software event]
dummy                                [Software event]
emulation-faults                     [Software event]
major-faults                         [Software event]
minor-faults                         [Software event]
page-faults OR faults                [Software event]
task-clock                           [Software event]

L1-dcache-load-misses                [Hardware cache event]
L1-dcache-loads                      [Hardware cache event]
```

# Perf Commands – perf stat

- Generate the statistics of the events that are occurred during process execution

```
$ perf stat ./prac_mutex
```

```
[jongbin@multicore-96:~/TA/Multicore/lab02$ perf stat ./prac_mutex
thread 140530927269632: local count -> 1000000
thread 140530910492416: local count -> 1000000
thread 140530902099712: local count -> 1000000
thread 140530893707008: local count -> 1000000
thread 140530885314304: local count -> 1000000
thread 140530876921600: local count -> 1000000
thread 140530868528896: local count -> 1000000
thread 140530860136192: local count -> 1000000
thread 140530851743488: local count -> 1000000
thread 140530843350784: local count -> 1000000
global count -> 10000000

Performance counter stats for './prac_mutex':

      15981.957836      task-clock:u (msec)      #    8.397 CPUs utilized
           0          context-switches:u      #    0.000 K/sec
           0          cpu-migrations:u        #    0.000 K/sec
          126          page-faults:u          #    0.008 K/sec
    3,999,649,184      cycles:u                #    0.250 GHz
    660,814,873      instructions:u            #    0.17 insn per cycle
    171,686,540      branches:u              #   10.743 M/sec
     3,501,680      branch-misses:u          #    2.04% of all branches

      1.903217149 seconds time elapsed
```

# Perf Commands – perf record

- Sampling events on per-thread, per-process, per-cpu basis and generate the output file

```
$ sudo perf record -g ./prac_mutex
```

```
thread 140498458892032: local count -> 1000000
thread 140498450499328: local count -> 1000000
thread 140498442106624: local count -> 1000000
thread 140498433713920: local count -> 1000000
thread 140498425321216: local count -> 1000000
thread 140498416928512: local count -> 1000000
thread 140498408535808: local count -> 1000000
thread 140498400143104: local count -> 1000000
thread 140498391750400: local count -> 1000000
thread 140498383357696: local count -> 1000000
global count -> 10000000
[ perf record: Woken up 18 times to write data ]
[ perf record: Captured and wrote 6.830 MB perf.data (57555 samples) ]
jongbin@multicore-96:~/TA/Multicore/lab02$ ls
perf.data  prac_mutex  prac_mutex.cpp
```

# Perf Commands – perf record

- Sampling events on a running process

```
$ sudo perf record -g -p 54487
```

```
jongbin@multicore-96:~/TA/Multicore/lab02$ ps u
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
jongbin	53728	0.0	0.0	22916	5600	pts/0	Ss	01:12	0:00	-bash
jongbin	54465	0.5	0.0	22880	5608	pts/4	Ss	03:33	0:00	-bash
jongbin	54487	0.0	0.0	98812	9580	pts/0	Sl+	03:33	0:05	./prac_mutex
jongbin	54498	0.0	0.0	47428	7512	pts/4	R+	03:33	0:00	ps u

```
$ sudo perf record -g -p `pidof prac_mutex`
```



# Perf Commands – perf report

- Analyze the output file generated from perf record

```
$ sudo perf report -g graph --no-children
```

```
Samples: 57K of event 'cycles:ppp', Event count (approx.): 715903980698041
```

	Overhead	Command	Shared Object	Symbol
+	30.00%	prac_mutex	[kernel.vmlinux]	[k] native_queued_spin_lock_slowpath
+	17.82%	prac_mutex	[kernel.vmlinux]	[k] futex_wake
+	8.91%	prac_mutex	libpthread-2.23.so	[.] pthread_mutex_unlock
+	8.60%	prac_mutex	[kernel.vmlinux]	[k] sys_futex
+	8.26%	prac_mutex	[kernel.vmlinux]	[k] try_to_wake_up
+	5.84%	prac_mutex	libpthread-2.23.so	[.] pthread_mutex_lock
+	5.53%	prac_mutex	[kernel.vmlinux]	[k] __unqueue_futex
+	4.30%	prac_mutex	[kernel.vmlinux]	[k] futex_wait_setup
+	3.99%	prac_mutex	[kernel.vmlinux]	[k] update_numa_stats
+	3.38%	prac_mutex	[kernel.vmlinux]	[k] __perf_event_task_sched_in
+	3.07%	prac_mutex	[kernel.vmlinux]	[k] update_cfs_rq_h_load
	0.31%	prac_mutex	[kernel.vmlinux]	[k] hash_futex
	0.00%	prac_mutex	[kernel.vmlinux]	[k] get_futex_value_locked
	0.00%	prac_mutex	[kernel.vmlinux]	[k] _raw_spin_lock
	0.00%	prac_mutex	libpthread-2.23.so	[.] __lll_lock_wait
	0.00%	prac_mutex	prac_mutex	[.] thread_func

# Perf Commands – perf annotate

- Analyze the output file generated from perf record in the instruction level

```
$ sudo perf annotate pthread_mutex_lock
```

```
pthread_mutex_lock /lib/x86_64-linux-gnu/libpthread-2.23.so
0.07 58: and $0x80,%edx
2.02 mov $0x1,%edi
0.03 xor %eax,%eax
0.03 mov %edx,%esi
40.92 lock cmpxchg %edi,(%r8)
      ↓ je 84
0.79 lea (%r8),%rdi
sub $0x80,%rsp
0.40 → callq __lll_lock_wait
0.63 add $0x80,%rsp
6.42 84: mov 0x8(%r8),%edx
test %edx,%edx
      ↓ jne 209
0.30 90: mov %fs:0x2d0,%eax
0.13 mov %eax,0x8(%r8)
2.12 addl $0x1,0xc(%r8)
```

# Perf Commands – perf annotate

- Analyze the output file generated from perf record in the instruction level

You can run the *annotate* command simply in the report screen

Samples: 57K of event 'cycles:ppp', Event count (approx.): 715903980698041

Overhead	Command	Shared Object	Symbol
+ 30.00%	prac_mutex	[kernel.vmlinux]	[k] native_queued_spin_lock_slowpath
+ 17.82%	prac_mutex	[kernel.vmlinux]	[k] futex_wake
- 8.91%	prac_mutex	libpthread-2.23.so	[.] pthread_mutex_unlock
pthread_mutex_unlock			
start_thread			
+ 8.60%	prac_mutex	[kernel.vmlinux]	[k] sys_futex
+ 8.26%	prac_mutex	[kernel.vmlinux]	[k] try_to_wake_up
+ 5.84%	prac_mutex	libpthread-2.23.so	[.] pthread_mutex_lock
+ 5.53%	prac_mutex	[kernel.vmlinux]	[k] __unqueue_futex
+ 4.30%	prac_mutex	[kernel.vmlinux]	[k] futex_wait_setup
+ 3.99%	prac_mutex	[kernel.vmlinux]	[k] update_numa_stats
+ 3.38%	prac_mutex	[kernel.vmlinux]	[k] __perf_event_task_sched_in
+ 3.07%	prac_mutex	[kernel.vmlinux]	[k] update_cfs_rq_h_load
0.31%	prac_mutex	[kernel.vmlinux]	[k] hash_futex
0.00%	prac_mutex	[kernel.vmlinux]	[k] get_futex_value_locked
0.00%	prac_mutex	[kernel.vmlinux]	[k] _raw_spin_lock
0.00%	prac_mutex	libpthread-2.23.so	[.] __lll_lock_wait
0.00%	prac_mutex	prac_mutex	[.] thread_func



pthread\_mutex\_unlock /lib/x86\_64-linux-gnu/libpthread-2.23.so

Disassembly of section .text:

0000000000000b4c0 <\_\_pthread\_mutex\_unlock>:  
\_\_GI\_\_pthread\_mutex\_unlock():

15.09	mov	0x10(%rdi),%esi
	mov	%rdi,%rdx
2.13	mov	%esi,%eax
0.27	mov	%esi,%r8d
0.11	and	\$0x17f,%eax
	and	\$0x7c,%r8d
	↓ jne	50
2.35	test	%eax,%eax
	↓ jne	60
0.21	1a: movl	\$0x0,0x8(%rdx)
0.53	subl	\$0x1,0xc(%rdx)
4.32	and	\$0x80,%esi
72.05	lock decl	(%rdx)
	↓ je	46
1.12	lea	(%rdx),%rdi

Choose a symbol and press 'a' key

# More about Perf...

---

Perf Tutorial

<https://perf.wiki.kernel.org/index.php/Tutorial>

# Thank You

---