

# Simple Two-phase locking with Readers-writer Lock

---

## Concurrent Programming

### Programming Project #2

Final due date: Oct 25, 2019 (HARD DEADLINE)

## 1 TASK OVERVIEW

Two-phase locking is one of widely-used concurrency control algorithms in database systems. This assignment is to **implement two-phase locking with readers-writer lock** within a simple structure.

## 2 TASK DETAILS

You have  $N$  worker threads and  $R$  records in your database. Each record is a 64 bit signed integer (initial value: 100). Each thread executes transactions until the *global execution or-*

*der* (initial value: 0) reaches the value  $E$ .

Each transaction consists of a sequence of following actions below.

1. Randomly pick up three different records  $i, j, k$ , respectively.
2. Acquire a global mutex that protects a lock table.
3. Acquire a reader lock for reading a value of the record  $i, R_i$ .
  - If it need to wait for acquiring the lock, do *deadlock checking*.
4. Release the global mutex.
5. Read  $R_i$ .
6. Acquire the global mutex again.
7. Acquire a writer lock for writing a value of the record  $j, R_j$ .
  - If it need to wait for acquiring the lock, do *deadlock checking*.
8. Release the global mutex.
9. Increase the value of  $R_j$  by 1, i.e.,  $R_j = R_j + R_i + 1$ .  
( $R_i$  is the value you have read at step 5)
10. Acquire the global mutex again.
11. Acquire a writer lock for writing a value of the record  $k, R_k$ .
  - If it need to wait for acquiring the lock, do *deadlock checking*.
12. Release the global mutex.
13. Decrease  $R_k$  by  $R_i$ . i.e.,  $R_k = R_k - R_i$ .  
( $R_i$  is the value you have read at step 5)  
  
———— committing phase ————
14. Acquire the global mutex.
15. Release all reader/writer locks acquired by this transaction.
16. Increase the global execution order by 1, and then fetch it as *commit\_id*. Initial value of the global execution order is 0 so that the first *commit\_id* need to be 1.

- If *commit\_id* is bigger than *E*, rollback all changes made by this transaction (Undo), and then release the global mutex, and terminate the thread.


Append a **commit log** into the thread#.txt with the below format:

*[commit\_id] [i] [j] [k] [R<sub>i</sub>] [R<sub>j</sub>] [R<sub>k</sub>]*

17. Release the global mutex. \_\_\_\_\_ committed \_\_\_\_\_

### 3 READERS-WRITER LOCK

Whenever a transaction accesses a record, it must acquire the corresponding record-level lock. You need to implement reader-writer lock for allowing multiple transactions to read the same record at the same time. Lock objects for a single record should be managed in a linked-list structure so that you can keep the order of the transactions accessing the record. When a transaction releases the lock, it must wake up other waiting transaction(s) who are safe to be proceed, if exists.




$\langle S_t$ : A reader lock of transaction  $t$ ,  $X_t$ : A writer lock of transaction  $t \rangle$


In the figure above, transaction 1 acquired three locks, and did not commit yet. Transaction 2 is acquiring a read lock for record 1, and waiting for acquiring a write lock for record 3. Transaction 3 is waiting for acquiring a read lock for record 2.

## 4 DEADLOCK CHECKING

Whenever a transaction needs to wait for acquiring a reader/writer lock, it must perform a deadlock checking. It can detect a dead lock by detecting a *wait-for* cycle in the lock table. If it turns out to be a deadlock, the transaction must rollback all changes it made (Undo), and restart.



< Deadlock example 1: transaction 1 -> transaction 2 -> transaction 1 >



< Deadlock example 2: transaction -> transaction 3 -> transaction 2 -> transaction 1 >

## 5 TEST PROTOCOL

Your program should take three command line arguments. Example of the execution is:

```
$ ./run N R E
```

*N*: Number of worker threads

*R*: Number of records

*E*: Last global execution order which threads will be used to decide termination

After your program finished,  $N$  output files need to be generated. The name of each output file is thread#.txt (i.e., if  $N$  is 3, thread1.txt, thread2.txt, thread3.txt need to be generated.) Each commit log should be printed on each line of the corresponding file.

The format of the output file is below: ( $N=3$ ,  $R=3$ )

thread1.txt
1 1 2 3 100 201 0
2 2 3 1 201 202 -101
4 3 1 2 102 2 200
...

thread2.txt
3 1 3 2 -101 102 302
...

thread3.txt
5 2 1 3 200 203 -98
...

## 6 SUBMISSION

You should upload your project into **project2** directory of your “hconnect” repository. You must create a proper **Makefile** in the project2 directory. The name of your executable file must be **run**, and it must be at the project2 directory as well.

You also need to upload your assignment report in your gitlab wiki page of your hconnect project. Please set a name of the wiki page as project2.

**Please enjoy this project and have fun !!**